

## HW1 Part 2: Explore `compromise_search.h`

Wenyi Liu

**Q: What's the big-O of `std::find`?**

A: `std::find` takes  $O(N)$  time because it is a linear search.

**Q: What is the big-O of `std::binary_search`?**

A: `std::binary_search` takes  $O(\log N)$  time. In binary search, we narrow down the search space by half. The search space reduces to 0 after  $1 + \log_2 N$  times and the while loop terminates.

**Q: What is the big-O of your code?**

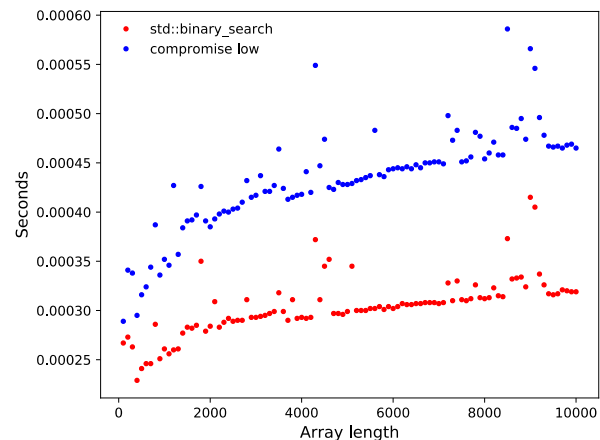
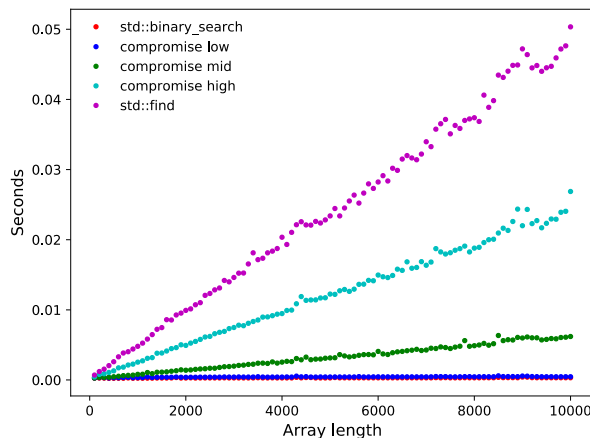
A: Without loss of generality, let  $N$  be the size of the initial search space, and  $n$  be the `small_size`. Then we need to do binary search until the search space reduce to  $(N - n)$ . This part takes  $O(\log(N-n))$  time. The remaining linear search takes  $O(n)$  time. So the overall time complexity of my code is  $O(\log(N-n) + n)$ .

Depending on the input, our code will behave differently.

If  $N \gg n$ , then the code takes  $O(\log N)$  time.

If  $n$  is comparable to  $N$ , then the code takes  $O(N)$  time.

The following two questions are based on the following figures. In the `test_code_performance.cpp`, I profile the running time of my code, `std::find` and `std::binary_search`. Compromise low, mid, high correspond to different `small_size`, with low being a fixed value of 10, mid being 20% of  $N$ , and high being 80% of  $N$ . And I use `performance.py` to make these two figures.



**Q: On your computer in practice, does your code appear to have the big-O you expect?**

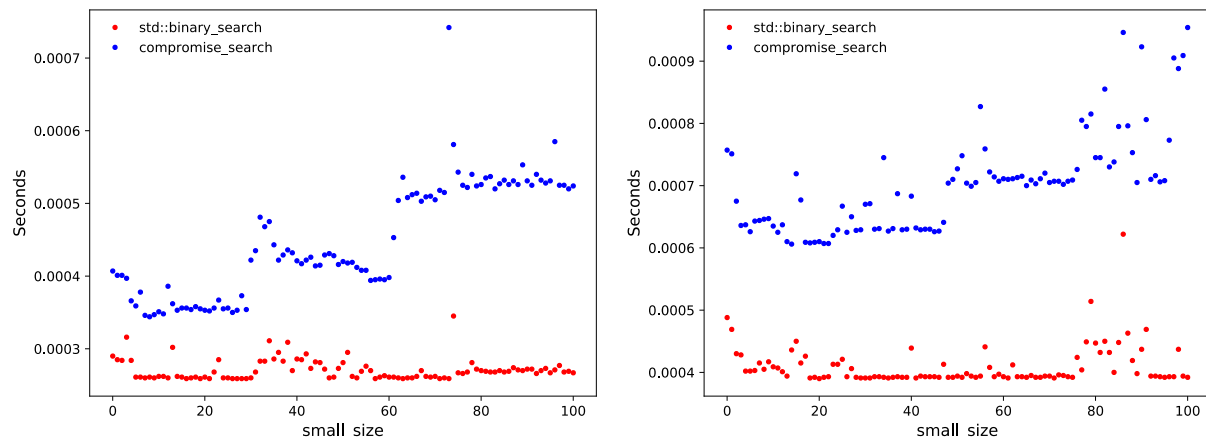
A: From above, the code should take  $O(\log(N-n) + n)$  time. So it should be slower than  $O(\log N)$  but faster than  $O(N)$ , which can

be seen from the left panel. The compromise search algorithm runs slower than `std::binary_search` but faster than `std::find`. When `small_size` is fixed at a very low value of 10, the compromise search is slightly slower than binary search. In this regime,  $N \gg n$  and compromise search should behave as  $O(\log N)$ , as can be seen from the right panel. When `small_size` is mid or high, i.e.  $n$  is comparable to  $N$ , the run time becomes  $O(N)$ , but still faster than `std::find`. Thus, the code has the expected big-O.

**Q: How fast is your code compared to `std::binary_search` in practice on your computer?**

A: From the right panel of the figure, compromised search is slightly slower than `std::binary_search` when the `small_size` is set to a fixed small constant. But both take  $O(\log N)$  time.

The following question uses the following figure. In the `choose_small_size.cpp`, I profile the running time of my code and `std::binary_search` v.s. `small_size`. I use `small_size.py` to make these two figures. Left panel has a search space of 1,000 and right panel has a search space of 100,000.



**Q: What's a good value for `small_size`?**

A: A good choice of `small_size` seems to consistently fall into the range between 10-30 regardless of the size of search space. At this size, `std::find` is as fast as `std::binarysearch`.